

Assignment 1: ROS Services and Launch Files

What to do

- Create a node named ‘array_server’ with a ROS2 service that accepts an array of fixed length (length of 8 elements) as its request.
- The service should return the average of the requested array as its response.
- Create a client node named ‘array_client’ that requests the average of the following array from your server: [10, 0, 5, -20, 13, 7, -9, -6]
- You should create a custom srv file for your communication between the client and the server.
- You should create two launch files to launch the client and the server nodes.

What to submit

Submit your src folder in a zip file. In addition, you should include screenshots of your client and servers running on two terminals simultaneously.

Assignment 2: Motion detection

What to do

1. You will have the video file intersection.mp4 in the videos folder. Split the video into images at 5 frames per second. Use the Linux program '`ffmpeg`'. Let's call this the dataset.
2. Write a Python file to process the generated images. In your Python file, read the first image from the dataset as background $B[0]$. Convert $B[0]$ to grayscale.
3. Write a loop to read all the other images in grayscale one by one. Each image you read will be $X[k]$. Apply the following equation.

$$B[k + 1] = B[k] + c(X[k] - B[k]), \quad (1)$$

where $B[k]$ is the current background image, $B[k + 1]$ is the updated background, $X[k]$ is the new image you read in the loop, and $(X[k] - B[k])$ is the difference between the current background image and the image you read.

4. The monadic operation $c(\cdot)$ is described below:

$$c(x) = \begin{cases} \sigma, & x > \sigma \\ x, & -\sigma \leq x \leq \sigma \\ -\sigma, & x < -\sigma \end{cases}$$

5. Use Python `numpy` library's `clip` function to easily apply the above function to your image $(X[k] - B[k])$.
6. Inside your loop, 1. At every iteration, save the image $(X[k] - B[k])$ in a separate folder, 2. concatenate the images $(X[k] - B[k])$ at every iteration in a variable and call it motion M , where

$$M[k + 1] = M[k] + (X[k] - B[k]).$$

7. $M[0]$ can be an empty image with only white-color pixels.
8. After the end of your loop, save both the updated background and the motion variable as two separate images.
9. Use `ffmpeg` to concatenate your individual motion images into a video.
10. (optional) Develop a colorized version of the motion image.

What to submit

Submit the background image, the motion image, the motion video, and Python code.

Assignment 4: Implementing CNN

Padding

Padding adds extra border pixels (typically zeros) around the input map before convolution, preventing edge information loss and controlling output size.

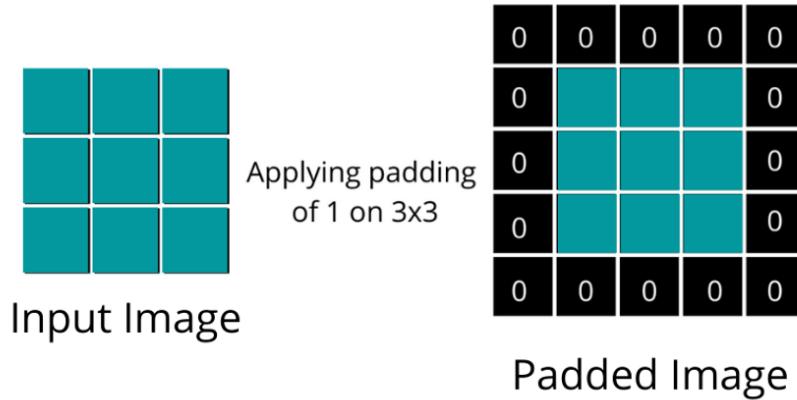


Figure 1: Padding

Strides

In convolutional neural networks, strides define the step size by which the filter (kernel) moves across the input feature map during convolution.

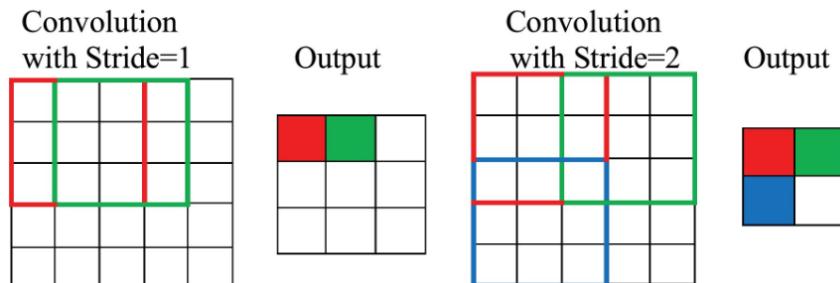


Figure 2: Strides

The following formula lets you map the output to the input with strides and padding.

$$\text{Output Size} = \left\lfloor \frac{\text{Input Size} + 2 \times \text{Padding} - \text{Kernel Size}}{\text{Stride}} \right\rfloor + 1$$

- Input Size: The height or width of the input feature map.
- Padding: The number of zero-padding layers added to the input borders.
- Kernel Size: The height or width of the convolutional filter.
- Stride: The step size by which the filter moves.
- $\lfloor \cdot \rfloor$: The floor function, rounding down to the nearest integer.

Implement the following CNN in Figure 3 using Tensorflow. Make the following assumption:

- Number of the classes at the output layer are 10.
- Assume the activation function is 'ReLU' for all the layers, but 'softmax' for the output layer.
- Assume the pooling layers are max pooling with 3x3 layers.

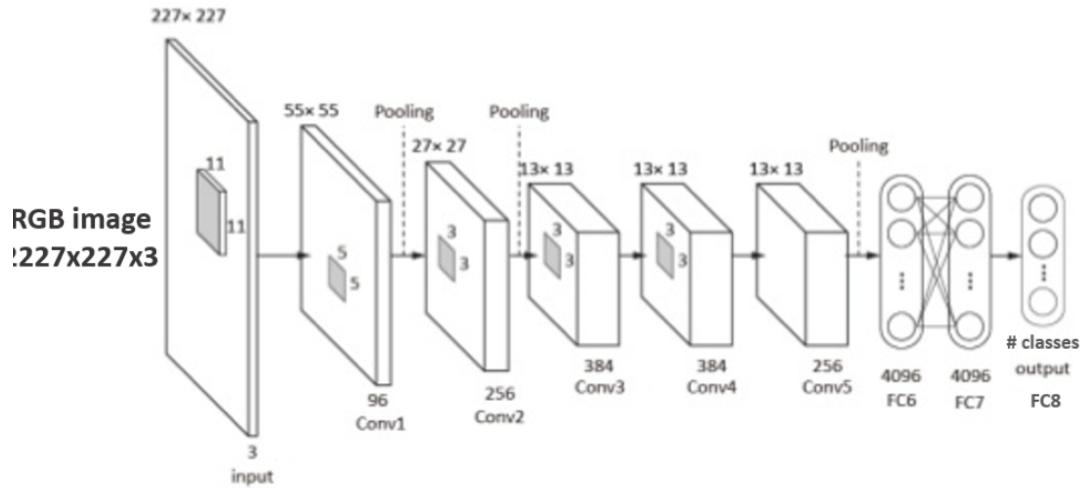


Figure 3: CNN

Final Project

Objective: At the end of the final project, you will have successfully implemented a YOLO object detection model on a mobile robot that controls via Raspberry Pi.

Goal: You are a member of the special police in the NYC. Successfully navigate the robot through the maze to identify the bad guys.

Training Steps

- Segment the video and tag them to their respective classes. Construct your Train set, Validation set, and Test set. I prefer this simple program, but there may be better options. **You are not allowed to use the first 2 minutes of the video for training or testing.**
- You will tag two objects : vehicles and bad-guys
- Train a YOLO by Ultralytics on the tagged images.
- Extract the weights in an appropriate format.

Implementation Steps

- Construct the GoPiGo robot and attach the Raspberry Pi. Connect the camera module to the Raspberry Pi.
- You can burn a base Ubuntu image on the RPi, connect the RPi to the internet, install ROS2 through this, and GoPiGo controls through this. **The alternative way would give you control over the libraries you install, hence it would help you debugging if needed.**
- Install YOLO through Ultralytics.
- Copy your trained weights to the RPi for inference.
- Write a Python script to read the serial images from the camera. Broadcast them via a ROS2 topic / service.
- Write another Python script to implement the YOLO inference. That script should also subscribe to the ROS2 topic / service to receive the images.
- Improve your Python scripts to recognize AprilTags and calculate the distance between the robot and the Tag.
- A third Python script performs GoPiGo controls depending on the results provided by the inference procedure.

Progress Tracking

- You have the option to work either by yourself or in groups of two. This is the signup link.
- You should have a GitHub repository to host your code. Your progress and contributions are tracked via github. If you decide to keep your repository private, you should grant me (chandima-ccsu) access to your repository.

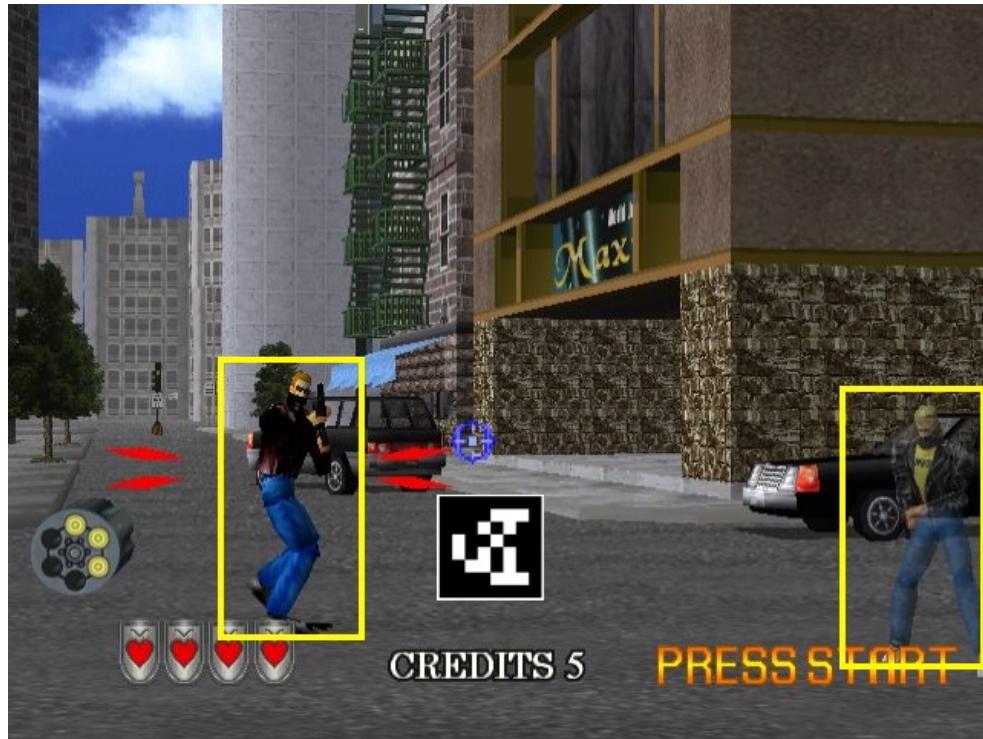


Figure 4: Inference time image with an AprilTag.

Grading (subject to adjustments)

Total available points: 50.

- Successfully Train a YOLO model (10pts).
- Implement the YOLO model on the RPi and correctly perform inference on test images (10pts).
- Successfully implement ROS2 Nodes for image communication and message passing on RPi (10pts).
- Successfully control the GoPiGo robot through the RPi (5pts).
- Successfully reading the AprilTags (5pts).
- A GitHub repository with a reasonable commit history, a README.md file with results of each of the above steps, and instructions to operate the robot (5pts).
- Working demonstration video (5pts).